

# Machine Learning Algorithms

## Algorithm

Algorithm	Page Number
1. k-means	1 - 3
2. k-nearest-neighbors	4 - 5
3. Linear Regression	6 - 7
4. Logistic Regression	8 - 9
5. SVM	10 - 11
6. Naive Bayes	12 - 13
7. PCA	14
8. Shallow Neural Network	15 - 22
9. k-means (short version)	23 - 25
10 Gradient Descent	26 - 29

## ns Algorithm

1

```
import numpy as np.  
  
def euclidean_distance(x1, x2):  
    return np.sqrt(np.sum((x1 - x2) ** 2))  
  
class KMeans:  
    def __init__(self, k=3, max_iters=100):  
        self.K = k  
        self.max_iters = max_iters  
        # list of sample indices for each cluster  
        self.clusters = [[] for _ in range(self.K)]  
        # mean feature vector for each cluster  
        self.centroids = []  
  
    def predict(self, X):  
        self.X = X  
        self.n_samples, self.n_features = X.shape  
        # initialize centroids  
        random_sample_idxs = np.random.choice(self.n_samples, self.K, replace=False)  
        self.centroids = [self.X[idx] for idx in random_sample_idxs]
```

(2)

# optimization

for i in range (self.max\_iters):

# update clusters

self.clusters = self.\_create\_clusters (self.centroids)

# update centroids

centroids-old = self.centroids

self.centroids = self.\_get\_centroids (self.clusters)

# check if converged

if self.\_is\_converged (centroids-old, self.centroids):

break

# return cluster labels

return self.\_get\_cluster\_labels (self.clusters)

def \_get\_cluster\_labels (self, clusters):

labels = np.empty (self.n\_samples)

for cluster\_idx, cluster in enumerate (clusters):

for sample\_idx in cluster:

labels [sample\_idx] = cluster\_idx

return labels

```
def _create_clusters(self, centroids):
    clusters = [[] for _ in range(self.K)]
    for idx, sample in enumerate(self.X):
        centroid_idx = self._closest_centroid(sample, centroids)
        clusters[centroid_idx].append(idx)
    return clusters
```

```
def -closest-centroid(self, sample, centroids):
    distances = [euclidean-distance(sample, point) for point in centroids]
    closest-idx = np.argmin(distances)
    return closest-idx
```

```
def -get-centroids (self, clusters):  
    centroids = np.zeros ((self.K, self.n_features))  
    for cluster_idx, cluster in enumerate (clusters):  
        cluster_mean = np.mean (self.X[cluster], axis=0)  
        centroids [cluster_idx] = cluster_mean  
  
    return centroids
```

```
def -is-converged (self, centroids-old, centroids):
    distances = [euclidean-distance (centroids-old[i], centroids[i])
                  for i in range (self.k)]
    return sum (distances) == 0.
```

(4)

## k-nearest neighbors Algorithm

import math.

# Euclidean distance calculation

```
def euclideanDistance (instance1, instance2, length):
    distance = 0
    for x in range (length):
        distance += math.pow (instance1[x] - instance2[x], 2)
    return math.sqrt (distance)
```

# Neighbors : selecting subset with the smallest distance

```
def getNeighbors (trainingSet, testInstance, k):
    distances = []
    length = len (testInstance) - 1
    for x in range (len (trainingSet)):
        dist = euclideanDistance (testInstance, trainingSet [x], length)
        distances.append ((trainingSet [x], dist))
    distances.sort (key=operator.itemgetter (1))
    neighbors = []
    for x in range (k):
        neighbors.append (distances [x][0])
    return neighbors.
```

(5)

KNN Algorithm

1. Load the data
2. Initialize K to the chosen number of neighbors
3. For each example in the data
  - (a) Calculate the distance b/w the query example and the current example from the data
  - (b) Add the distance and the index of the example to an ordered collection.

4. Sort the ordered collection of distances and indices from smallest to largest (in ascending order) by the distances

5. Pick the first K entries from the sorted collection

6. Get the labels of the selected K entries

7. If regression, return the mean of the K labels

8. If classification, return the mode of the K labels.

def getResponse(neighbors):

    classVotes = {}

    for x in range(len(neighbors)):

        response = neighbors[x][-1]

        if response in classVotes:

            classVotes[response] += 1

    else:

        classVotes[response] = 1

sortedVotes = sorted(classVotes.items(), key=operator.itemgetter(1), reverse=True)

return sortedVotes[0][0]

def getAccuracy(testSet, predictions):

    correct = 0

    for x in range(len(testSet)):

        if testSet[x][-1] in predictions[x]:

            correct += 1

return (correct / float(len(testSet))) \* 100

def main():

    predictions = []

    k = 5

    for x in range(len(testSet)):

        neighbors = getNeighbors(trainingSet, testSet[x], k)

        result = getResponse(neighbors)

        predictions.append(result)

accuracy = getAccuracy(testSet, predictions)

# Linear Regression Algorithm

```
import numpy as np
def mean-squared-error(y_true, y_pred):
    return np.mean((y_true - y_pred) ** 2)
```

class LinearRegression : fit (selected) (a)

```
def __init__(self, learning_rate=0.001, n_iters=1000):
```

self. Ir = learning-rate

$$\text{self}.\text{n\_iters} = \text{n\_iters}$$

self. weights = None

self . bias = None

```
def fit(self, X, y):
```

• n-samples, n-features = X, shape

# initialize parameters

self. weights = np.zeros (n-features)

$$\text{self-bias} = 0.$$

# gradient descent

```
for _ in range(self.n_iters):
```

`y-predicted = np.dot(X, self.weights) + self.bias`

# compute gradients

$$dw = (1/n\_samples) * np.dot(x.T, (y\_predicted - y))$$

$$db = (1/n\text{-samples}) * np.sum(y\text{-predicted} - y)$$

# update parameters

self-weights = self-for \* dw

$$\text{self\_bias} = \text{self\_lr} * \text{db}$$

(7)

def predict(self, x):

$$y\text{-approximated} = \text{np}\cdot\text{dot}(x, \text{self}\cdot\text{weights}) + \text{self}\cdot\text{bias}$$

return y-approximated.

regressor = LinearRegression(learning\_rate=0.01, n-iterations=1000)

regressor.fit(x-train, y-train)

predictions = regressor.predict(x-test)

mse = mean-squared-error(y-test, predictions)

(8)

## Logistic Regression Algorithm

```
import numpy as np
```

```
def accuracy(y_true,  
accuracy = np.sum(
```

```
class LogisticRegression:
```

```
def __init__(self, lr=0.001, epochs=1000):
```

```
    self.lr = lr
```

```
    self.epochs = epochs
```

```
    return accuracy
```

```
def sigmoid(self, z):
```

```
    return 1 / (1 + np.exp(-z))
```

```
def cost_function(self, X, y, weights):
```

```
    z = np.dot(X, weights)
```

```
    predict_1 = y * np.log(self.sigmoid(z))
```

```
    predict_0 = (1 - y) * np.log(1 - self.sigmoid(z))
```

```
    return -np.sum(predict_1 + predict_0) / X.shape[0]
```

```
def fit(self, X, y):
```

```
    loss = []
```

```
    weights = np.random.rand(X.shape[1])
```

```
    for _ in range(self.epochs):
```

```
        y_hat = self.sigmoid(np.dot(X, weights))
```

```
        weights -= self.lr * np.dot(X.T, y_hat - y) / X.shape[0]
```

```
        loss.append(self.cost_function(X, y, weights))
```

⑨

```
def predict(self, x):
    z = np.dot(x, self.weights)
    return np.array([1 if i > 0.5 else 0 for i in self.sigmoid(z)])
```

  

```
regressor = LogisticRegression()
regressor.fit(X-train, y-train)
predictions = regressor.predict(X-test)
accuracy(y-test, predictions)
```

10

## SVM Algorithm

```
import numpy as np
```

```
class SVM:
```

```
    def __init__(self, learning_rate=0.001, lambda_param=0.01, n_iters=100):
```

```
        self.lr = learning_rate
```

```
        self.lambda_param = lambda_param
```

```
        self.n_iters = n_iters
```

```
        self.w = None
```

```
        self.b = None
```

```
    def fit(self, X, y):
```

```
        y_ = np.where(y <= 0, -1, 1)
```

```
        n_samples, n_features = X.shape
```

```
        self.w = np.zeros(n_features)
```

```
        self.b = 0.
```

```
        for _ in range(self.n_iters):
```

```
            for idx, x_i in enumerate(X):
```

```
                condition = y_[idx] * (np.dot(x_i, self.w) - self.b) >= 1
```

```
                if condition:
```

```
                    self.w -= self.lr * (2 * self.lambda_param * self.w)
```

```
                else:
```

```
                    self.w -= self.lr * (2 * self.lambda_param * self.w -
```

```
                    self.b -= self.lr * y_[idx]
```

```
                    np.dot(x_i, y_[idx]))
```

~~self.weights, loss~~

11

def predict(self, x):

linear\_output = np.dot(x, self.w) - self.b

return (np.sign(linear\_output))

clf = SVM()

clf.fit(x, y)

predictions = clf.predict(x)

def pif(clf, x0, x1):

numerator = self.w[clf.predict(x0)] \* np.dot(x1 - x0, x1)

denominator = np.sqrt((x0 - x1).dot(x0 - x1)) \* np.sqrt((x1 - x0).dot(x1 - x0))

cosine = numerator / denominator

angle = np.arccos(cosine) \* 180 / np.pi

return angle

def pif2(clf, x0, x1):

numerator = self.w[clf.predict(x0)] \* np.dot(x1 - x0, x1)

denominator = np.sqrt((x0 - x1).dot(x0 - x1)) \* np.sqrt((x1 - x0).dot(x1 - x0))

## Naive Bayes Algorithm

(12)

class NaiveBayes:

def fit(self, X, y):

n-samples, n-features = X.shape

self.\_classes = np.unique(y)

n-classes = len(self.\_classes)

# initialize mean, variance, priors

self.\_mean = np.zeros((n-classes, n-features), dtype=np.float64)

self.\_var = np.zeros((n-classes, n-features), dtype=np.float64)

self.\_priors = np.zeros(n-classes, dtype=np.float64)

for c in self.\_classes:

X\_c = X[c == y]

self.\_mean[c, :] = X\_c.mean(axis=0)

self.\_var[c, :] = X\_c.var(axis=0)

self.\_priors[c] = X\_c.shape[0] / float(n-samples)

def predict(self, X):

y-pred = [self.\_predict(x) for x in X]

return (y-pred)

def -predict (self, x):

postiors = []

for idx, c in enumerate (self.\_classes):

prior = np.log (self.\_priors [idx])

class-conditional = np.sum (np.log (self.\_pdf (idx, x)))

posterior = prior + class-conditional

postiors.append (posterior)

return (self.\_classes [np.argmax (postiors)])

def -pdf (self, class\_idx, x):

mean = self.\_mean [class\_idx]

var = self.\_var [class\_idx]

numerator = np.exp (-((x-mean)\*\*2 / (2\*var)))

denominator = np.sqrt (2 \* np.pi \* var)

return (numerator / denominator)

(14)

## PCA Algorithm

class PCA:

```
def __init__(self, n_components):
    self.n_components = n_components
    self.components = None
    self.mean = None
```

def fit(self, X):

```
    self.mean = np.mean(X, axis=0)
```

```
X = X - self.mean
```

```
cov = np.cov(X.T)
```

```
eigenvalues, eigenvectors = np.linalg.eig(cov)
```

```
eigenvectors = eigenvectors.T
```

```
idxs = np.argsort(eigenvalues)[-1:]
```

```
eigenvalues = eigenvalues[idxs]
```

```
eigenvectors = eigenvectors[idxs]
```

```
self.components = eigenvectors[0:self.n_components] # store first n components
```

def transform(self, X):

# project data

```
X = X - self.mean
```

```
return (np.dot(X, self.components.T))
```

## Neural Network

```
at numpy as np
in sklearn.metrics import accuracy_score.
```

```
def relu(z):
    return np.maximum(z, 0)
```

```
def relu_prime(z):
    return (z > 0).astype(z.dtype)
```

```
def sigmoid(z):
    return 1 / (1 + np.power(np.e, -z))
```

```
def sigmoid_prime(z):
    return z * (1 - z)
```

```
def leaky_relu(z, alpha = 0.01):
    return np.where(z > 0, z, z * alpha)
```

```
def leaky_relu_prime(z, alpha = 0.01):
    dz = np.ones_like(z)
    dz[z < 0] = alpha
    return dz
```

```
def tanh(z):
    return np.tanh(z)
```

```
def tanh_prime(z):
    return 1 - (tanh(z) ** 2)
```

10

```
def get-activation-function(name):  
    if name == 'relu':  
        return relu  
    elif name == 'sigmoid':  
        return sigmoid  
    elif name == 'leaky-relu':  
        return leaky_relu  
    elif name == 'tanh':  
        return tanh  
    else:  
        raise ValueError('Only "relu", "leaky-relu", "tanh" and "sigmoid" supported')
```

```
def get-derivative-activation-function(name):  
    if name == 'relu':  
        return relu_prime  
    elif name == 'sigmoid':  
        return sigmoid_prime  
    elif name == 'leaky-relu':  
        return leaky_relu_prime  
    elif name == 'tanh':  
        return tanh_prime  
    else:  
        raise ValueError('Only "relu", "leaky-relu", "tanh" and "sigmoid" supported')
```

(17)

initialize\_nn(  $X, y$ , hidden-nodes = 4, random-state = None) :

if random-state != None :

    np.random.seed(random-state)

n-x = X.shape[1]

n-h = hidden-nodes

n-y = y.shape[1]

w1 = np.random.randn(n-x, n-h) # np.sqrt(1./n-h)

b1 = np.zeros((1, n-h))

w2 = np.random.randn(n-h, n-y) # np.sqrt(1./n-y)

b2 = np.zeros((1, n-y))

params = {

'w1': w1,

'b1': b1,

'w2': w2,

'b2': b2

}

return params

(18)

def forward-prop (x, nn-params, activation = 'relu'):

    W1 = nn-params['W1']

    b1 = nn-params['b1']

    W2 = nn-params['W2']

    b2 = nn-params['b2']

    Z1 = np.dot(x, W1) + b1

    A1 = get-activation-function(activation)(Z1)

    Z2 = np.dot(A1, W2) + b2

    A2 = get-activation-function('sigmoid')(Z2)

    result = {

        'Z1': Z1,

        'A1': A1,

        'Z2': Z2,

        'A2': A2

}

return result

calculate - loss ( $A_2, y$ ):

```
m = y.shape[0]
return np.squeeze(-(1./m) * np.sum(np.multiply(y, np.log(A2)) +
np.multiply(np.log(1-A2), 1-y)))
```

# squeeze will convert [[cost]] to 'cost' float variable

```
def backward-prop (x, y, m-params, cache, learning-rate = 0.001,
activation = 'relu'):
```

$w_1 = m\text{-params}['w1']$

$b_1 = m\text{-params}['b1']$

$w_2 = m\text{-params}['w2']$

$b_2 = m\text{-params}['b2']$

$z_1 = \text{cache}['z1']$

$a_1 = \text{cache}['a1']$

$z_2 = \text{cache}['z2']$

$a_2 = \text{cache}['a2']$

$m = x.shape[0]$

-  $dz_2 = A^2 - y$

$dw_2 = (1./m) * \text{np.dot}(a1.T, dz_2)$

$db_2 = (1./m) * \text{np.sum}(dz_2, \text{axis}=0, \text{keepdims=True})$

20

$dz1 = np.dot(dz2, w2.T) * \text{get-derivative-activation-function}(\text{activation})(A1)$ .

$$dwl = (1/m) * np.dot(x.T, dz)$$

```
dbl = (1./m) * np.sum(dz1, axis=0, keepdims=True)
```

update = {

'dwl': dwl,

'dbl': dbl,

'dw<sup>2</sup>': dw<sup>2</sup>,

'db2': db<sup>2</sup>

}

return updates

```
def update-weights (m-params, updates, learning-rate = 0.01) :
```

$w_1 = m - \text{param} [w_1]$

$$|b| = \text{m-param} [^c b ^l]$$

$w_2 = m - \text{params} [w_2]$

$$b^2 = m - \text{params} [{}'b^2']$$

`dwl = updates['dwl']`

$db1 = \text{updates} [^{'db1'}]$

$d\mathbf{w}^2 = \text{updates} [d\mathbf{w}^2]$

$db^2 = \text{update} [db^2]$

$$W1 = W1 - \text{learning-rate} * dW1$$

$$b1 = b1 - \text{learning-rate} * db1$$

$$W2 = W2 - \text{learning-rate} * dW2$$

$$b2 = b2 - \text{learning-rate} * db2$$

final-result = {

'W1': W1,

'b1': b1,

'W2': W2,

'b2': b2

}

return final-result.

def predict(parameters, X, y):

cache = forward-prop(X, parameters, 'relu')

predictions = cache['A2'] > 0.5

return predictions.

(22)

```
n = 10000
learning-rate = 0.01
hidden-nodes = 50
nn-params = initialize-nn (x, y, hidden-nodes = hidden-nodes, random-state = 0)
history = []
for i in range (epoch+1):
    result = forward-prop (x-train, nn-params, 'relu')
    loss = calculate-loss (result ['A2'], y-train)
    history.append (loss)
    update = backward-prop (x-train, y-train, nn-params, result, learning-rate=0.01,
                            activation='relu')
    nn-params = update-weights (nn-params, update, learning-rate)
if i% (epoch/10) == 0:
    print ('Epoch: {} \t Loss: {:.6f} \t Train Accuracy: {:.3f}'.format (i, loss,
                                                                     accuracy-score (y-train, predict (nn-params, x-train, y-train)),
                                                                     accuracy-score (y-test, predict (nn-params, x-test, y-test))))
```

## K Means

Outline:

Assume we have input data points  $x_1, x_2, x_3, \dots, x_n$  and the value of  $K$  (the number of clusters). We follow the procedure:

1. Pick  $K$  points as the initial centroids from the dataset, either randomly or the first  $K$ .
2. find the euclidean distance of each point in the dataset with the identified  $K$  points (cluster centroids).
3. Assign each data point to the closest centroid using the distance calculated in the previous step.
4. Find the new centroid by taking the average of the points in each cluster group.
5. Repeat 2 to 4 for a fixed number of iterations or till the centroids don't change.

class:

def \_\_init\_\_(self, k=3, max\_iterations=500):

self.k = k

self.max\_iterations = max\_iterations

def euclidean\_distance(self, point1, point2):

# return math.sqrt((point1[0] - point2[0])\*\*2 + (point1[1] - point2[1])\*\*2)

return np.linalg.norm(point1 - point2, axis=0)

def fit(self, data):

# let the first K points from the dataset be the initial centroids

self.centroids = {}

for i in range(self.k):

self.centroids[i] = data[i]

# start k-means clustering.

for i in range(self.max\_iterations):

# create classes the size of k

self.classes = {}

for j in range(self.k):

self.classes[j] = []

# find the distance between the points and the centroids

for point in data:

distances = []

for index in self.centroids:

distances.append(self.euclidean\_distance(point, self.centroids[index]))

# start k-means clustering.

for i in range (self.max\_iterations):

    # create cluster the size of k

    self.classes = {}

    for j in range (self.k):

        self.classes[j] = []

    # find the distance between the points and the centroids.

    for point in data:

        distances = []

        for index in self.centroids:

            distances.append (self.euclidean\_distance (point, self.centroids[index]))

        # find which cluster the datapoint belongs to by finding the min distance.

        cluster\_index = distances.index (min (distances))

        self.classes[cluster\_index].append (point)

    # now that we have classified the datapoints into clusters,

    # we need to again find the new centroid by taking the

    # average of the points in the cluster class.

    for cluster\_index in self.classes:

        self.centroids[cluster\_index] = np.average (self.classes[cluster\_index], axis=0)

cost function for linear regression model

$$\text{MSE}(x, \theta_0) = \frac{1}{m} \sum_{i=1}^m (\theta^T x^{(i)} - y^{(i)})^2$$

m training examples:  $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$ .  
 $x^{(i)} \in \mathbb{R}^{n_x \times 1}$

$$X = \begin{bmatrix} 1 & 1 & & & 1 \\ x^{(1)} & x^{(2)} & \dots & & x^{(m)} \\ 1 & 1 & & & 1 \end{bmatrix} \begin{matrix} \uparrow \\ n_x \\ \downarrow \\ \end{matrix}$$

← m →

$$X \in \mathbb{R}^{n_x \times m}$$

$$y = [y^{(1)} \ y^{(2)} \ \dots \ y^{(m)}] \in \mathbb{R}^{1 \times m}$$

$$\theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_{n_x} \end{bmatrix}$$

To find the value of  $\theta$  that minimizes the cost function,  
 there is a closed form solution.

$$\hat{\theta} = (X^T X)^{-1} X^T y$$

## Batch Gradient Descent

Partial derivatives of the cost function,

$$\frac{\partial}{\partial \theta_j} \text{MSE}(\theta) = \frac{2}{m} \sum_{i=1}^m (\theta^T x^{(i)} - y^{(i)}) x_j^{(i)}$$

Gradient vector of the cost function

$$\nabla_{\theta} \text{MSE}(\theta) = \begin{pmatrix} \frac{\partial}{\partial \theta_0} \text{MSE}(\theta) \\ \frac{\partial}{\partial \theta_1} \text{MSE}(\theta) \\ \vdots \\ \frac{\partial}{\partial \theta_n} \text{MSE}(\theta) \end{pmatrix}$$

$$= \frac{2}{m} X^T (X \cdot \theta - y)$$

Gradient Descent step

$$\theta^{(\text{next step})} = \theta - \eta \nabla_{\theta} \text{MSE}(\theta).$$

eta = 0.1 # learning rate.

n\_iterations = 1000

m = 100

theta = np.random.randn(2, 1).

for iteration in range(n\_iterations):

gradients = 2/m \* X.T.dot(X.dot(theta) - y)

theta = theta - eta \* gradients.

## No Gradient Descent

(28)

⑥ epochs = 50.  
t<sub>0</sub>, t<sub>1</sub> = 5, 50 # learning schedule hyperparameters.

ef learning-schedule (t):  
return t<sub>0</sub> / (t + t<sub>1</sub>)

theta = np.random.randn(2, 1) # random initialization

for epoch in range(n-epochs):

for i in range(m):

random-index = np.random.randint(m)

x<sub>i</sub> = X[:, random-index : random-index + 1]

y<sub>i</sub> = y[random-index : random-index + 1]

gradients = 2 \* x<sub>i</sub>.T. dot (x<sub>i</sub> . dot (theta) - y<sub>i</sub>)

eta = learning-schedule (epoch \* m + i)

theta = theta - eta \* gradients.

## Mini-Batch Gradient Descent

(29)

```
def get-batches (x, y, batch-size, i):  
    x-new = x [i:i+batch-size, :]  
    y-new = y [i:i+batch-size]  
    return x-new, y-new
```

```
def MiniBatchGradientDescent (x, y, theta, learning-rate = 0.1, batch-size = 10):  
    num-batches = int (x.shape[0] / batch-size)  
  
    gradient = 0  
  
    for i in range (0, num-batches):  
        x-batch, y-batch = get-batches (x, y, batch-size, i)  
        hypothesis = np.dot (x-batch, theta)  
        loss = hypothesis - y-batch  
        x-tran = x-batch.T  
        gradient = np.dot (x-tran, loss) / batch-size  
        theta = theta - learning-rate * gradient  
  
    return theta
```